

# Machine Learning: from Theory to Practice

Lecture 5: Trees and ensemble methods

F. d'Alché-Buc and E. Le Pennec  
`florence.dalche@telecom-paristech.fr`

Fall 2016

- 1 Reminder
- 2 Decision and regression trees
- 3 Ensemble methods
  - Bagging
  - Random forests
  - AdaBoost as a Greedy Scheme
  - Gradient Boosting
- 4 References

## Supervised statistical learning

- $\mathcal{S}_n = \{(x_1, y_1), \dots, (x_n, y_n)\}$  is an i.i.d sample of size  $n$ , drawn from the joint probability law  $P(X, Y)$  fixed but unknown.
- $\mathcal{H}$  a class of functions and a local loss function  $\ell : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$

Solve the following optimization problem:

$$\min_{h \in \mathcal{H}} R(h)$$

with  $R(h) = \mathbb{E}[\ell(h(x), y)]$

Example: Prediction loss,  $\ell(h(x), y) = 0$  if  $h(x) = y$  and 1 otherwise

In practise : minimize  $R_n(h)$ , the empirical loss, while controlling the complexity of  $h$

## Methodology

- Define
  - a representation space for data
  - a class of functions (a class of hypotheses) where to find the solution
  - a loss function to be minimized
  - an optimization algorithm
  - a model selection method for hyperparameters

# Today's agenda

Reminder


- Decision and Regression Trees
- Ensemble methods

- 1 Reminder
- 2 Decision and regression trees
- 3 Ensemble methods
  - Bagging
  - Random forests
  - AdaBoost as a Greedy Scheme
  - Gradient Boosting
- 4 References

# Decision trees

Decision and  
regression trees

Invented between 1979 and 1983 simultaneously by L. Breiman (CART, applied stats) et col. et R. Quinlan (ID3, a new discipline called Machine Learning)




## *Reference:*

Breiman, Leo; Friedman, J. H.; Olshen, R. A.; Stone, C. J. (1984). Classification and regression trees. Monterey, CA: Wadsworth Brooks/Cole Advanced Books Software.

# Decision trees 1


Decision and  
regression trees





# Decision trees 1

Decision and  
regression trees




## First idea:

Use not one but many linear separators to build non linear decision frontiers

## Second idea

# Decision trees 1

Decision and  
regression trees



## First idea:


Use not one but many linear separators to build non linear decision frontiers

## Second idea

Use hyperplane of the form  $x^j = c$  to keep an interpretation of the learned function

# Decision tree 2

Decision and  
regression trees




## Third idea

The learned function can be represented by a tree structure whose each node is associated to a hyperplane  $x^j = \theta_j$  and each leaf to a constant function i.e. a class index.

# Decision tree 3

Decision and  
regression trees



At the end of the learning phase you know every feature involved in the decision function

The tree encodes rules of logic  $0^+$

if  $(x^{j_1} > c_{j_1})$  and  $(x^{j_2} \leq c_{j_2})$  and ... then  $x$  belongs to class  $k$

# Linear separator orthogonal to one vector for the canonical basis

Decision and  
regression trees

$x^j$  : continuous variable

$$t_{j,c}(\mathbf{x}) = \text{sign}(x^j - c) \quad (1)$$

$x^j$  : categorical variable here with 2 values  $\{v_1^j, \dots, v_2^j\}$ :

$$t(x; v_j) = \mathbb{I}(x^j = v_j) \quad (2)$$

# A new class of constant piece-wise functions

Decision and  
regression trees

We consider a new class of functions :

## Definition

$$\forall x \in \mathbb{R}^p, h^{tree}(x) = \sum_{j=1}^M \alpha_m \mathbb{I}_{\mathcal{C}_m}(x),$$

where  $\mathcal{C}_1 \cup \dots \mathcal{C}_M$  is a partition of  $\mathbb{R}^p$  and each  $\mathcal{C}_m$  is defined by a subset of hyperplanes orthogonal to the canonical basis.

Let  $x \in \mathbb{R}^p$ . Denote  $m(x) \in \{1, \dots, M\}$  such that  $x \in \mathcal{C}_{m(x)}$ .

- Classification, by majority vote in  $\mathcal{C}_{m(x)}$ :

$$h_n(x) = \alpha_{m(x)} = \arg \max_{k=1, \dots, K} \sum_{x_i \in \mathcal{C}_{m(x)}} \mathbb{I}(y_i = k)$$

- Regression, by empirical mean in  $\mathcal{C}_{m(x)}$ :

$$f_n(x) = \alpha_{m(x)} = \frac{1}{|\mathcal{C}_{m(x)}|} \sum_{x_i \in \mathcal{C}_{m(x)}} y_i$$

Goal: learn  $\mathcal{C}_1 \cup \dots \mathcal{C}_M$ .

# A local and greedy learning algorithm

Decision and  
regression trees

## (Binary tree)

- 1 Let  $\mathcal{S}$  be the current training set
- 2 Let us build a root node
- 3 Find the best split  $t(\mathbf{x})$  to apply to the current training set  $\mathcal{S}$  such that the local loss  $L(t, \mathcal{S})$  be minimal
- 4 Attach the chosen split to the current node and separate the current training dataset into two subsets  $\mathcal{S}_d$  et  $\mathcal{S}_g$  with the help of the split
- 5 Build a right node and a left node.
- 6 Measure the stopping criterion on the right side and if checked, then the right node becomes a leaf , otherwise go to 3 with  $\mathcal{S}_d$  as the current training subset
- 7 Measure the stopping criterion on the left side and if checked, then the left node becomes a leaf , otherwise go to 3 with  $\mathcal{S}_g$  as the current training subset



Let  $\mathcal{S}$  be the current training set and  $t_{j,\tau}$  a binary split.

Note

- $\mathcal{D}(\mathcal{S}, j, \tau) = \{(\mathbf{x}, y) \in \mathcal{S}, t_{j,\tau}(\mathbf{x}) > 0\}$
- $\mathcal{G}(\mathcal{S}, j, \tau) = \{(\mathbf{x}, y) \in \mathcal{S}, t_{j,\tau}(\mathbf{x}) \leq 0\}.$

$\tau_1, \dots, \tau_C$  can be chosen regularly spaced or using an histogram.

Among all the parameters  $(j, \tau) \in \{1, \dots, p\} \times \{\tau_1, \dots, \tau_C\}$ , we search for  $\hat{j}$  and  $\hat{\tau}$  that minimize:

## Local loss function

$$L(t_{j,\tau}, \mathcal{S}) = \frac{n_d}{n} H(\mathcal{D}(\mathcal{S}, j, \tau)) + \frac{n_g}{n} H(\mathcal{G}(\mathcal{S}, j, \tau))$$

$$n_d = |\mathcal{D}(\mathcal{S}, j, \tau)|$$

$$n_g = |\mathcal{G}(\mathcal{S}, j, \tau)|$$

where  $H(\mathcal{S})$  measures the impurity of a subset  $\mathcal{S}$

# Impurity criteria for supervised classification 1/2

Decision and  
regression trees

For a given dataset  $\mathcal{S}$  of  $n$  labeled data and for each class  $k$ :

$$p_k(\mathcal{S}) = \frac{1}{n} \sum_{i=1}^n \mathbb{I}(y_i = k)$$

here are the main criteria  $H$  that can be used:

**Cross entropy**

$$H(\mathcal{S}) = - \sum_{k=1}^K p_k(\mathcal{S}) \log p_k(\mathcal{S})$$

# Impurity criteria for supervised classification 2/2

Decision and  
regression trees

## Cross entropy

$$H(\mathcal{S}) = - \sum_{k=1}^K p_k(\mathcal{S}) \log p_k(\mathcal{S})$$

## Gini index

$$H(\mathcal{S}) = \sum_{k=1}^K p_k(\mathcal{S})(1 - p_k(\mathcal{S}))$$


## Classification error

$$H(\mathcal{S}) = 1 - p_{C(\mathcal{S})},$$

with  $C(\mathcal{S})$ : majority class in  $\mathcal{S}$ .

# Visualization of impurity criteria

Decision and  
regression trees



Stop the construction if one of the following value is reached

- Maximal Depth
- Maximal number of leaves
- Minimal number of training data in a leaf

NB : If the minimal number of training data in a leaf is equal to 1, it is possible that the tree grows until a perfect classification of the training examples is achieved: **overfitting** !

## Some remarks: categorical variables, multi-class classification

Decision and  
regression trees

- To get a binary tree, if a categorical variable can take  $K$  values, it can be transformed into  $K$  binary variables.
- The learning algorithm is appropriate to solve binary classification problem as well as multi-class classification problems

We maximize the variance reduction !

$$L(t_{j,\tau}, \mathcal{S}) = \text{VAR}_{emp}(\mathcal{S}) - \frac{n_d}{n} \text{VAR}_{emp}(\mathcal{D}(j, \tau, \mathcal{S})) - \frac{n_g}{n} \text{VAR}_{emp}(\mathcal{G}(j, \tau, \mathcal{S}))$$

$$\text{VAR}_{emp}(\mathcal{S}) = \frac{1}{|\mathcal{S}|} \sum_{(x_i, y_i) \in \mathcal{S}} (y_i - \bar{y})^2$$

We attempt to get homogeneous outputs !



(1) We focus on the selection of on of the following hyperparameters:

- Maximal Depth
- Maximal number of leaves
- Minimal number of training data in a leaf

typically by → **cross-validation**.

## (2) **by pruning**

Another dataset, called validation set, is used to re-visit a tree that was grown until covering all the training examples. One only keeps branches that bring a performance improvement on validation set.

# Advantages and drawbacks

Decision and  
regression trees

## Advantages

- Nonlinear interpretable decision frontier → explainable AI
- Consistency (see Scott and Nowak, IEEE Trans. Inf. Theory, 2006)
- Very efficient computation of images of  $h_n^{tree}$  or  $f_n^{tree}$  : complexity in time in (*depth*).
- Works for continuous and categorical features
- Works directly for multi-class classification
- Very flexible

## Drawbacks

- Large variance Estimator : a change in the root implies a change in the whole tree
- Learning a tree is a NP-complete problem (Hyafil and Rivest, Information Proc. Letters, 1976)
- No global learning algorithm

# Extension of tree approaches to other problem

Decision and  
regression trees

Like  $k$ -nearest neighbors, SVM/SVR, trees can be adapted to other machine learning tasks:

- classification with imbalanced datasets
- multiple output prediction
- quantile regression
- ranking
- anomaly detection
- time series modeling

How ?

- Change the local loss (impurity criteria)

- 1 Reminder
- 2 Decision and regression trees
- 3 Ensemble methods**
  - Bagging
  - Random forests
  - AdaBoost as a Greedy Scheme
  - Gradient Boosting
- 4 References

# Ensemble methods for classification and regression

Ensemble methods

- Machine Learning not so "automatic": too many hyperparameters to tune
- ① **Committee learning** or **wisdom of the crowd**: better results are obtained by combining the predictions of a set of **diverse** classifiers/regressors
- ② **Meta-learning**: a procedure to automatically use a base classifier/regressor even weak to produce a performant classifier/regressor
- ③ **Ensemble learning**: is a kind of meta-learning, improves upon a single predictor by building an ensemble of predictors (with no hyperparameter)

- 1995: Boosting, Freund and Schapire
- 1996: Bagging, Breiman
- 1999: GradientBoosting, Friedman et al.
- 2001: Random forests, Breiman
- 2006: Extra-trees, Geurts, Ernst, Wehenkel

- Improve upon a single predictor by building an ensemble of predictors (with no hyperparameter)
- → meta-learning: the parameter of the *meta-learning* algorithm is those of the base learner and the size of the ensemble



Let  $f_t, t = 1, \dots, T$  be  $T$  different regressors.

Notations:

$$\epsilon_t(x) = y - f_t(x)$$

$$MSE(f_t) = \mathbb{E}[\epsilon_t(x)^2]$$

$$\begin{aligned} f_{ens}(x) &= \frac{1}{T} \sum_t f_t(x) \\ &= y - \frac{1}{T} \sum_t \epsilon_t(x). \end{aligned}$$

## Encourage the diversity of base predictors

Ensemble methods

$$MSE(f_{ens}) = \mathbb{E}[(y - f_{ens}(x))^2] = \frac{1}{T^2} \mathbb{E}[(\sum_t \epsilon_t(x))^2]$$

Now, if  $\epsilon_t$  are mutually independent with zero mean, then we have:

$$MSE(f_{ens}) = \frac{1}{T^2} \mathbb{E}[\sum_t \epsilon_t(x)^2]$$

The more diverse are the classifiers, the more we reduce the mean square error !

- **Encourage the diversity of base predictors by:**
  - using bootstrap samples (Bagging and Random forests)
  - using randomized predictors (ex: Random forests)
  - using weighted version of the current sample (Boosting) with weights dependent on the previous predictor (adaptive sampling)

- 1 Reminder
- 2 Decision and regression trees
- 3 Ensemble methods**
  - Bagging
    - Random forests
    - AdaBoost as a Greedy Scheme
    - Gradient Boosting
- 4 References

# Decomposition bias/variance in regression

Ensemble methods

Given  $x$ ,

$$E_S[E_{y|x}[(y - f_S(x))^2]] = \text{noise}(x) + \text{bias}^2(x) + \text{variance}(x) \quad (3)$$

$\text{noise}(x)$ :  $E_{y|x}[(y - E[y|x])^2]$ :

quantifies the error made by the Bayes model ( $E[y|x]$ )

$$\text{bias}^2(x) = (E[y|x] - E_S[f_S(x)])^2$$

measures the difference between minimal error (Bayes error) and the average model

$$\text{variance}(x) = E_S[(f_S(x) - E_S[f_S(x)])^2]$$

measures how much  $f_S(x)$  varies from one training set to another

Assume we can generate several training samples  $\mathcal{S}_1, \dots, \mathcal{S}_T$  from  $P(x, y)$ .

A first algorithm:

- draw  $T$  training samples  $\{\mathcal{S}_1, \dots, \mathcal{S}_T\}$
- learn a model  $f_t \in \mathcal{F}$  from each training sample  $\mathcal{S}_t; t = 1, \dots, T$
- compute the average model :  $f_{ens}(x) = \frac{1}{T} \sum_{t=1}^T f_t(x)$

The bias remains the same:

$$\text{bias}(x) = E_{S_1, \dots, S_T}[f_{\text{ens}}(x)] = \frac{1}{T} \sum_t E_{S_t}[f_t(x)] = E_S[f_t(x)]$$

The variance is divided by T:

$$E_{S_1, \dots, S_T}[(f_{\text{ens}}(x) - E_{S_1, \dots, S_T}[f_{\text{ens}}(x)])^2] = \frac{1}{T} E_S[(f_S(x) - E_S[f_S(x)])^2]$$

# Bagging (Breiman 1996)

In practice, we do not know  $P(x,y)$  and we have only one training sample  $\mathcal{S}$ .

Bagging = Bootstrap Aggregating:

- draw  $T$  bootstrap samples  $\{\mathcal{B}_1 \dots, \mathcal{B}_T\}$  from  $\mathcal{S}$
- Learn a model  $f_t$  for each  $\mathcal{B}_t$
- Build the average model:  $f_{bag}(x) = \frac{1}{T} \sum_t f_t(x)$



Efron 1977.

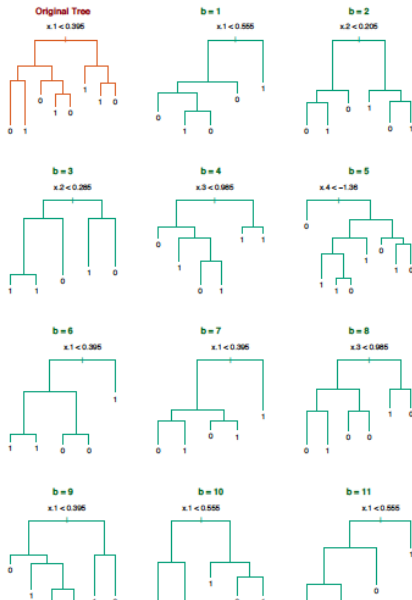
- Bootstrap is a robust estimation method
- A bootstrap sample drawn from  $\mathcal{S}$  is a sample of the same size obtained by uniformly drawing examples from  $\mathcal{S}$  with replacing
- If we construct  $B$  bootstrap samples and apply an estimation procedure  $T$  on  $\mathcal{S}_1, \dots, \mathcal{S}_B$  then we can compute:
  - A bootstrap average:  $\bar{T} = \frac{1}{B} \sum_{b=1}^B T(\mathcal{S}_b)$
  - A bootstrap variance :  $VAR_{boot}(T) = \frac{1}{B} \sum_{b=1}^B (T(\mathcal{S}_b) - \bar{T})^2$

The bootstrap variance is often used to choose among several hyperparameteres which gives the most stable estimator

# Example of bagged trees

Ensemble methods

[Book: The elements of statistical learning, Hastie, Tibshirani,



- Variance is reduced but the bias can increase a bit (the effective size of a bootstrap sample is 30% smaller than the original training set  $\mathcal{S}$ )
- The obtained model is however more complex than a single model
- Bagging works for unstable predictors (neural nets, trees)
- In supervised classification, bagging a good classifier usually makes it better but bagging a bad classifier can make it worse

- 1 Reminder
- 2 Decision and regression trees
- 3 Ensemble methods**
  - Bagging
  - **Random forests**
  - AdaBoost as a Greedy Scheme
  - Gradient Boosting
- 4 References

- Perturbe and combine algorithms
  - Perturbe the base predictor
  - Combine the perturbed predictors

REFS: Random forests: Breiman 2001

Geurts, Ernst, Wehenkel, Extra-trees, 2006

## Random forests algorithm

- INPUT: candidate feature  $F$ ,  $\mathcal{S}_{train}$
- for  $t=1$  to  $T$ 
  - $\mathcal{B}_{train}^{(t)} \leftarrow$  a bootstrap sample from  $\mathcal{S}_{train}$
  - $h_{tree}^{(t)} \leftarrow$  randomized decision tree learned from  $\mathcal{B}_{train}^{(t)}$
- OUTPUT:  $H^T = \frac{1}{T} \sum_t h_{tree}^{(t)}$

- To select a split at a node:
  - $R_f(F) \leftarrow$  randomly select (without replacement)  $f$  feature splits from  $F$  with  $f \ll |F|$
  - Choose the best split in  $R_f(F)$  (consider the different cut-points)
- Do not prune this tree

## Learning a single randomized tree:

- To select a split at a node:
  - $R_K(F) \leftarrow$  randomly select (without replacement)  $f$  feature splits from  $F$  with  $f \ll |F|$
  - Choose the best split in  $R_f(F)$  (consider the different cut-points)
- Do not prune this tree



## Extra-trees

- INPUT: candidate feature splits  $F = \{1, \dots, p\}, \mathcal{S}_{train}$
- for  $t=1$  to  $T$ 
  - Always use  $\mathcal{S}_{train}$
  - $h_{tree}^{(t)}$ : extremely randomized decision tree learned from  $\mathcal{S}_{train}$
- OUTPUT:  $H^T = \frac{1}{T} \sum_t h_{tree}^{(t)}$


## Learning a single randomized tree in extra-trees:

- To select a split at a node:
  - randomly select (without replacement)  $K$  feature splits from  $F$  with  $K \ll |F|$
  - Draw  $K$  splits using the procedure  $\text{Pick-a-random-split}(\mathcal{S}, i)$ :
    - let  $a_{max}^i$  and  $a_{min}^i$  denote the maximal and minimal value of  $x_i$  in  $\mathcal{S}$
    - Draw uniformly a cut-point  $a_c$  in  $[a_{max}^i, a_{min}^i]$
- Choose the best split among the  $K$  previous splits

Do not prune this tree

# Random forest


Example of decision frontier:



# Comparison (just an example)

Ensemble methods

[Book: The elements of statistical learning, Hastie, Tibshirani, Friedman, 2001]



- 1 Reminder
- 2 Decision and regression trees
- 3 Ensemble methods**
  - Bagging
  - Random forests
  - AdaBoost as a Greedy Scheme**
  - Gradient Boosting
- 4 References

- **Is it possible to "boost" a weak learner into a strong learner ?** Michael Kearns
- Yoav Freund and Rob Schapire proposed an iterative scheme, called, Adaboost to solve this problem
  - **Idea**: train a sequence of learners on weighted datasets with weights depending on the loss obtained so far.
  - Freund and Schapire received the Godel prize in 2003 for their work on AdaBoost.

# Boosting a linear classifier

Ensemble methods


$$H_1(x) = h_1(x)$$

$$\text{Binary Classifier: } F_1(x) = \text{sign}(H_1(x))$$


Here:  $h_1$ : linear classifier

Training error =  $R_n$

$t = 0$



$t = 1$




*Source Jiri Matas (Oxford U.)*

# Boosting a linear classifier


Ensemble methods

$$H_2(x) = \alpha_1 h_1(x) + \alpha_2 h_2(x)$$

Binary Classifier:  $F_2(x) = \text{sign}(H_2(x))$



$R_n(H_t)$




*Source Jiri Matas (Oxford U.)*




# Boosting a linear classifier

Ensemble methods



$R_n(H_t)$





t Source Jiri

Matas (Oxford U.)

# Boosting a linear classifier



Ensemble methods



*Source Jiri Matas (Oxford U.)*

# Boosting a linear classifier


Ensemble methods




Source Jiri Matas (Oxford U.)

# Boosting a linear classifier

Ensemble methods





$R_n(H_t)$



*Source Jiri Matas (Oxford U.)*

# Boosting a linear classifier

Ensemble methods



*Source Jiri Matas (Oxford U.)*

## Definition: weak classifier

A classifier whose average training error is no more than 0.5

NB : it means that we do not need to have a deep architecture as the base classifier (a "short" tree will fit for instance, a linear classifier will be perfect and so on...)

- ①  $\mathcal{H}$ : a chosen class of "weak" binary classifiers,  $\mathcal{A}$ : a learning algorithm for  $\mathcal{H}$
- Set  $w_1(i) = 1/n$ ;  $H_0 = 0$
- For  $t = 1$  to  $T$ 
  - $h_t = \arg \min_{h \in \mathcal{H}} \epsilon_t(h)$
  - with  $\epsilon_t(h) = \mathbb{P}_{i \sim \mathbf{w}_t}[h(x_i) \neq y_i]$
  - Choose  $\alpha_t$
  - Choose  $w_{t+1}$
  - $H_t = H_{t-1} + \alpha_t h_t$
- Output  $F_T = \text{sign}(H_T)$

$\mathcal{H}$ : a chosen class of "weak" binary classifiers

- Set  $w_1(i) = 1/n$ ;  $H_0 = 0$
- For  $t = 1$  to  $T$ 
  - $h_t = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n \epsilon_t(h)$
  - With  $\epsilon_t(h) = \mathbb{P}_{i \sim \mathbf{w}_t}[h(x_i) \neq y_i]$
  - $\epsilon_t = \epsilon_t(h_t)$
  - $\alpha_t = \frac{1}{2} \log \frac{1-\epsilon_t}{\epsilon_t}$
  - let  $w_{t+1,i} = \frac{w_t(i)e^{-\alpha_t y_i h_t(x_i)}}{Z_{t+1}}$  where  $Z_{t+1}$  is a renormalization constant such that  $\sum_{i=1}^n w_{t+1,i} = 1$
- $H_t = H_{t-1} + \alpha_t h_t$

Output  $F_T = \text{sign}(H_T)$



# What weight to choose ?

With the chosen definition, we have:

$$\begin{aligned}w_{t+1,i} &= \frac{w_{t,i} e^{-\alpha_t y_i h_t(x_i)}}{Z_t} \\&= \frac{w_{t-1,i} e^{-\alpha_{t-1} y_i h_{t-1}(x_i)} e^{-\alpha_t y_i h_t(x_i)}}{Z_{t-1} Z_t} \\&= \frac{e^{-y_i \sum_{s=1}^t \alpha_s h_s(x_i)}}{n \prod_{s=1}^t Z_s} \\&= \frac{e^{-y_i H_t(x_i)}}{n \prod_{s=1}^t Z_s}\end{aligned}$$

You see the weights encourage to correct examples badly classified by the whole combination  $H_t$

# First of all let us study $Z_t$

Ensemble methods

$$\begin{aligned}Z_t &= \sum_{i=1}^n w_t(i) e^{-\alpha_t y_i h_t(x_i)} \\&= \sum_{i=1}^n w_t(i) e^{-\alpha_t y_i h_t(x_i)} \\&= \sum_{i: y_i h_t(x_i)=+1} w_t(i) e^{-\alpha_t} + \sum_{i: y_i h_t(x_i)=-1} w_t(i) e^{\alpha_t} \\&= (1 - \epsilon_t) e^{-\alpha_t} + \epsilon_t e^{\alpha_t} \\&= (1 - \epsilon_t) \sqrt{\frac{\epsilon_t}{1 - \epsilon_t}} + \epsilon_t \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \\&= \dots \\&= 2\sqrt{\epsilon_t(1 - \epsilon_t)}\end{aligned}$$

## Theorem

The empirical error of the classifier returned by Adaboost at time  $T$  verifies:

$$R_n(F_T) \leq e^{-2 \sum_{t=1}^T (\frac{1}{2} - \epsilon_t)^2}.$$

Furthermore, if for all  $t \in [1, T]$ ,  $\gamma \leq (\frac{1}{2} - \epsilon_t)$ , then

$$R_n(F_T) \leq e^{-2\gamma^2 T}.$$

For all  $u \in \mathbb{R}$ , we have  $1_{u \leq 0} \leq \exp(-u)$ .

Then

$$\begin{aligned} R_n(F_T) &= \frac{1}{n} \sum_{i=1}^n 1_{y_i F_T(x_i) \leq 0} \\ &\leq \frac{1}{n} \sum_{i=1}^n \exp(-y_i F_T(x_i)) = \frac{1}{n} \sum_{i=1}^n \left[ n \prod_{t=1}^T Z_t \right] w_{t+1,i} = \prod_{t=1}^T Z_t \end{aligned}$$

We can now express  $\prod Z_t$  in terms of  $\epsilon_t$ :

$$\begin{aligned}\prod_{t=1}^T Z_t &= \prod_{t=1}^T 2\sqrt{\epsilon_t(1-\epsilon_t)} \\ &= \text{by remarkable identity} \\ &= \prod_{t=1}^T \sqrt{1-4(1/2-\epsilon_t)^2} \\ &\leq \prod_t e^{-2(1/2-\epsilon_t)^2} = e^{-2\sum_{t=1}^T (1/2-\epsilon_t)^2}\end{aligned}$$

using the identity  $1-u \leq \exp(-u)$ .

The proof reveals several interesting properties:

- ❶  $\alpha_t$  is chosen to minimize  $\prod_t Z_t = g(\alpha)$  with  $g(\alpha) = (1 - \epsilon_t)e^{-\alpha} + \epsilon_t e^{\alpha}$ 
  - $g'(\alpha) = -(1 - \epsilon_t)e^{-\alpha} + \epsilon_t e^{\alpha}$
  - $g'(\alpha) = 0$  iff  $(1 - \epsilon_t)e^{-\alpha} = \epsilon_t e^{\alpha}$  iff  $\alpha = 1/2 \log \frac{1-\epsilon_t}{\epsilon_t}$
- ❷ The equality  $(1 - \epsilon_t)e^{-\alpha} = \epsilon_t e^{\alpha}$  means that Adaboost assigns at each time  $t$  the same distribution mass to correctly classified examples and incorrectly classified ones. However there is no contradiction because the number of incorrectly examples decreases.

<http://scikit-learn.org/stable/modules/ensemble.html#adaboost>

```
>>> from sklearn.crossvalidation import crossva/score
>>> from sklearn.datasets import loadiris
>>> from sklearn.ensemble import AdaBoostClassifier
>>> iris = loadiris()
>>> clf = AdaBoostClassifier(nestimators=100)
>>> scores = crossvalscore(clf, iris.data, iris.target)
>>> scores.mean() 0.9...
```

- 1 Reminder
- 2 Decision and regression trees
- 3 Ensemble methods**
  - Bagging
  - Random forests
  - AdaBoost as a Greedy Scheme
  - Gradient Boosting**
- 4 References



At the same time, different groups proved that Adaboost writes as a coordinate descent in the convex hull of  $\mathcal{H}$ .

- Greedy function approximation, Friedman, 1999.
- MarginBoost and AnyBoost : Mason et al. 1999.

- At each boosting step, one need to solve

$$(h_t, \alpha_t) = \arg \min_{h, \alpha} \sum_{i=1}^n \ell(y_i, H_{t-1}(x_i) + \alpha h(x_i)) = L(y, H_{t-1} + \alpha h)$$

- Gradient approximation

$$L(y, H_{t-1} + \alpha h) \sim L(y, H_{t-1}) + \alpha \langle \nabla L(H_{t-1}), h \rangle.$$

- Gradient boosting: replace the minimization step by a *gradient descent* type step:
  - Choose  $h_t$  as the best possible descent direction in  $\mathcal{H}$
  - Choose  $\alpha_t$  that minimizes  $L(y, H + \alpha h_t)$
- Easy if finding the best descent direction is easy!

# Gradient boosting and adaboost

Those two algorithms are equivalent!

- Denoting  $H_t = \sum_{t'=1}^t \alpha_{t'} h_{t'}$ ,

$$\begin{aligned}\sum_{i=1}^n e^{-y_i(H_{t-1}(x_i) + \alpha h(x_i))} &= \sum_{i=1}^n e^{-y_i H_{t-1}(x_i)} e^{-\alpha y_i h(x_i)} \\ &= \sum_{i=1}^n w'_i(t) e^{-\alpha y_i h(x_i)} \\ &= (e^\alpha - e^{-\alpha}) \sum_{i=1}^n w'_i(t) \ell^{0/1}(y_i, h(x_i)) \\ &\quad + e^{-\alpha} \sum_{i=1}^n w'_i(t)\end{aligned}$$

- The minimizer  $h_t$  in  $h$  is independent of  $\alpha$  and is also the minimizer of

$$\sum_{i=1}^n w'_i(t) \ell^{0/1}(y_i, h(x_i))$$

- The optimal  $\alpha_t$  is then given by

$$\alpha_t = \frac{1}{2} \log \frac{1 - \epsilon'_t}{\epsilon'_t}$$

with  $\epsilon'_t = (\sum_{i=1}^n w'_i(t) \ell^{0/1}(y_i, h_t(x_i))) / (\sum_{i=1}^n w'_i(t))$

- One verify then by recursion that

$$w_i(t) = w'_i(t) / \left( \sum_{i=1}^n w'_i(t) \right)$$

and thus the two procedures are equivalent!

# AnyBoost or Forward Stagewise Additive model

Ensemble methods

- General greedy optimization strategy to obtain a linear combination of *weak* predictor
  - Set  $t = 0$  and  $H_0 = 0$ .
  - For  $t = 1$  to  $T$ ,
    - $(h_t, \alpha_t) = \arg \min_{h, \alpha} \sum_{i=1}^n \ell(y_i, H_{t-1}(x_i) + \alpha h(x_i))$
    - $H_t = H_{t-1} + \alpha_t h_t$
  - Output  $H_T = \sum_{t=1}^T \alpha_t h_t$

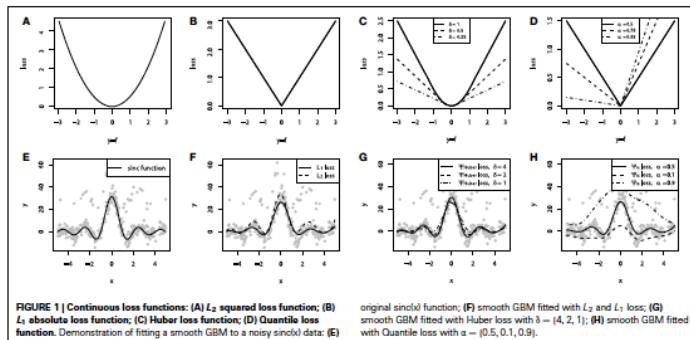
# Losses in Forward Stagewise Additive Modeling

Ensemble methods

- AdaBoost with  $\ell(y, h) = e^{-yh}$
- LogitBoost with  $\ell(y, h) = \log(1 + e^{-yh})$
- $L_2$ Boost with  $\ell(y, h) = (y - h)^2$  (Matching pursuit)
- $L_1$ Boost with  $\ell(y, h) = |y - h|$
- HuberBoost with
$$\ell(y, h) = |y - h|^2 \mathbf{1}_{|y-h| < \epsilon} + (2\epsilon|y - h| - \epsilon^2) \mathbf{1}_{|y-h| \geq \epsilon}$$

Simple principle but no easy numerical scheme except for AdaBoost and  $L_2$ Boost...


# Continuous loss functions and gradient boosting



- Loss function for regression:  $\ell(y, h) = (y - h)^2$
- $(h_t, \alpha_t) = \arg \min_{h, \alpha} \sum_{i=1}^n (y_i - H_t(x_i) + \alpha h(x_i))^2$

Fitting the residuals.





- You have to wait a long time to see Boosting overfit. However contrary to first assertions, Adaboost does overfit
- The surrogate loss which is empirically minimized is not exactly the prediction loss.
- Early stopping may be a first answer
- More interestingly : shrinkage

## *Reference:*

Boosting Algorithms: learning, regularization, prediction, Buhlman, Hothorn, Statistical science, 2007.

- 1 Reminder
- 2 Decision and regression trees
- 3 Ensemble methods
  - Bagging
  - Random forests
  - AdaBoost as a Greedy Scheme
  - Gradient Boosting
- 4 References

- Perrone, Cooper, When classifiers disagree, 1992
- Tumer and Gosh, 1996
- Breiman, Bagging predictors, 1996
- Buhlman and Yu, Analyzing bagging, Annals of stats., 2002
- Breiman, Random Forests, Machine Learning, 2001.
- Geurts, Ernst, Wehenekl, Extra-trees, JMLR, 2006
- Boosting:
  - Freund and Schapire, 1996
  - Greedy function approximation, Friedman, 1999.
  - MarginBoost and AnyBoost : Mason et al. 1999.
  - Tutorial Boosting de Buehlman
  - Boosting, Schapire and Freund, 2012